

Convolutional Neural Network using TensorFlow


Youngjae Yu (yj.yu@vision.snu.ac.kr)



SEOUL NATIONAL UNIV.
VISION & LEARNING


TensorFlow Basics

- Graph example



TensorFlow Basics

- Graph example




TensorFlow Basics

- Constant
 - `tf.constant(value, dtype, ...)`
 - e.g. `tf.constant([1, 2, 3])`
- Variable
 - `tf.Variable(initial-value, ...)`
 - e.g. `tf.Variable(tf.zeros(shape=(2,2)))`
- Placeholder
 - `tf.placeholder(dtype, shape, ...)`
 - e.g. `tf.placeholder(tf.float32, shape=(10, 10))`


Softmax Regression

- Multinomial logistic regression




Neural Network

- XOR Problem




Neural Network

- If only with linear classifier (e.g. logistic regression model), ...




Neural Network

- But if we combine them, ...




Neural Network

- Put hidden layer between input and output layer!
- Consists of fully-connected layers in which neurons between two adjacent layers are fully pairwise connected.
- Neurons within a single layer share no connections.




Neural Network


- Activation Functions




Sigmoid



tanh



ReLU




Leaky ReLU

Regularization


Regularization

- Let's make the model which explains the data below.
- $y = f(x) + \epsilon$



Regularization


- Three candidates



Need to choose one of them... Which one is the best?

Regularization


- Three candidates



Need to choose one of them... Which one is the best?

Regularization

- Overfitting problem



Regularization

- Several methods to avoid overfitting problem
 - L2 regularization
 - Dropout
 - ...


Regularization

- L2 regularization
 - Add L2 penalty (λw^2) to cost function
 - `tf.nn.l2_loss(t, ...)`


```
with graph.as_default():  
    ...  
  
    loss = tf.reduce_mean(  
        tf.nn.softmax_cross_entropy_with_logits(logits,  
            + l2_lambda * tf.nn.l2_loss(weights)  
    )  
    ...
```

Regularization

- Dropout
 - Sampling a Neural Network within the full Neural Network, and only updating the parameters of the sampled network based on the input data



(a) Standard Neural Net



(b) After applying dropout.


Regularization

- Dropout
 - `tf.nn.dropout(x, keep_prob, ...)`
 - `keep_prob`: The probability that each element is kept


```
with graph.as_default():  
    ...  
  
weights_1 = tf.Variable(tf.truncated_normal([image_size  
biases_1 = tf.Variable(tf.zeros([size_of_hidden])))  
logits_1 = tf.matmul(tf_train_dataset, weights_1) + bia  
output_1 = tf.nn.relu(logits_1)  
  
dropped_output_1 = tf.nn.dropout(x=output_1,  
                                keep_prob=keep_prob)  
  
    ...
```

Convolutional Neural Network

Convolution




Input I : $[W_I, H_I, D_I] \in \mathbb{R}^3$




Filter F : $[W_F, H_F, D_F, K] \in \mathbb{R}^4$

Convolution




Convolution




Convolution

- Stride




Convolution

- Stride




Convolution

- Padding




Convolution

- Padding




Convolution

- Padding




Convolution

- Padding



Convolution

- Padding



Convolution

- Padding

0	0	0	0
0	0	0	0
0	1	4	0
0	2	5	0
0	3	6	0
0	0	0	0
0	0	0	0

Zero


6	3	6	3
5	2	5	2
4	1	4	1
5	2	5	2
6	3	6	3
5	2	5	2
4	1	4	1

Reflect


2	2	5	5
1	1	4	4
1	1	4	4
2	2	5	5
3	3	6	6
3	3	6	6
2	2	5	5

Symmetric

Pooling




Input I : $[W_I, H_I, D_I] \in \mathbb{R}^3$



Filter F : $[W_F, H_F, D_F] \in \mathbb{R}^3$

Pooling



Max: $O_{x,y} = \max_{i,j} I_{x+i,y+j}$

Average: $O_{x,y} = \frac{1}{W_F \times H_F} \sum_i^{W_F} \sum_j^{H_F} I_{x+i,y+j}$

Reference

- TensorFlow official webpage

<https://www.tensorflow.org/>

- Stanford CS class - CS231n

<http://cs231n.github.io/>

- Udacity - Deep learning course

<https://www.udacity.com/course/deep-learning--ud730>